

1. Short Answer Questions

Q1: Explain the primary differences between TensorFlow and PyTorch. When would you choose one over the other?

Primary Differences:

- **Computational Graph:**
 - **TensorFlow:** Primarily uses a **static computational graph**. This means the graph is defined first and then executed. This can offer benefits in terms of optimization for deployment and distributed computing once the graph is finalized. However, it can make debugging more challenging as changes require recompiling the graph.
 - **PyTorch:** Uses a **dynamic computational graph** (also known as "define-by-run"). The graph is built on the fly as operations are performed. This makes PyTorch more "Pythonic" and significantly easier to debug, as you can inspect tensors and graph operations during runtime.
- **Ease of Use/Pythonic Nature:**
 - **PyTorch:** Generally considered more Pythonic and easier for beginners to pick up due to its imperative programming style, which aligns well with standard Python programming. Debugging is often likened to debugging regular Python code.
 - **TensorFlow:** While TensorFlow 2.x with Keras has greatly improved its user-friendliness and adopted eager execution (dynamic graph-like behavior), its underlying architecture can still feel more rigid and less intuitive for those accustomed to pure Python.
- **Production Readiness:**
 - **TensorFlow:** Historically, TensorFlow has had a stronger reputation and more mature tools for production deployment (e.g., TensorFlow Serving, TensorFlow Lite). Its static graph allows for more aggressive optimizations and deployment to various platforms.
 - **PyTorch:** While PyTorch's production deployment capabilities have significantly improved with tools like TorchScript, TensorFlow still often has an edge in large-scale, enterprise-level deployments, especially when aiming for extreme optimization.
- **Community and Ecosystem:**

- **TensorFlow:** Backed by Google, TensorFlow has a massive and established community with extensive resources, tutorials, and a vast ecosystem of related tools (TensorBoard, TF Serving, TF Lite, TF.js, etc.).
- **PyTorch:** While newer, PyTorch has a rapidly growing and highly active community, especially popular in academic research due to its flexibility. Its ecosystem is expanding rapidly.

When to choose one over the other:

- **Choose PyTorch when:**
 - You are a beginner or prefer a more "Pythonic" and intuitive experience.
 - You are engaged in research and rapid prototyping, where flexibility, dynamic graph manipulation, and easier debugging are crucial.
 - You prioritize quick experimentation and iterating on model architectures.
 - You are working on projects with less stringent deployment requirements initially.
- **Choose TensorFlow when:**
 - You are building large-scale, production-ready applications that require robust deployment, scalability, and integration with various platforms (mobile, web, edge devices).
 - You need extensive tooling for monitoring, serving, and managing models in a production environment (MLOps).
 - You benefit from a larger, more mature ecosystem with more pre-built models and distributed training capabilities out-of-the-box for certain scenarios.
 - You are comfortable with a more defined-and-then-run workflow, or using Keras as a high-level API.

Q2: Describe two use cases for Jupyter Notebooks in AI development.

1. **Exploratory Data Analysis (EDA) and Preprocessing:** Jupyter Notebooks are invaluable for the initial stages of AI development. Data scientists can load datasets, visualize distributions, check for missing values, perform feature engineering, and apply various preprocessing steps (like normalization or encoding) interactively. The ability to execute code cell by cell, immediately see the output (including plots and dataframes), and add explanatory Markdown text makes it easy to understand the data, identify patterns, and

prepare it for model training. This iterative process is crucial for gaining insights into the data's characteristics.

2. **Rapid Prototyping and Model Experimentation:** Jupyter Notebooks excel at quickly building, training, and evaluating different machine learning models. Researchers and developers can define model architectures, train them on small subsets of data, and test different hyperparameter configurations in an iterative fashion. The immediate feedback loop provided by cell execution allows for quick adjustments and comparisons of model performance, accelerating the experimentation process. You can easily modify a model, retrain it, and visualize its performance metrics or predictions without needing to re-run the entire script, making it ideal for trying out various ideas efficiently.

Q3: How does spaCy enhance NLP tasks compared to basic Python string operations?

spaCy provides a powerful and efficient framework for Natural Language Processing (NLP) that goes far beyond basic Python string operations in several key ways:

1. **Linguistic Annotation and Structure:** Basic Python string operations (like `split()`, `find()`, `replace()`) treat text as a sequence of characters. spaCy, on the other hand, understands and annotates text with rich linguistic information. It performs:
 - **Tokenization:** Breaks text into meaningful units (tokens) while handling complex cases like contractions, punctuation, and hyphenated words.
 - **Part-of-Speech (POS) Tagging:** Identifies the grammatical role of each word (e.g., noun, verb, adjective).
 - **Dependency Parsing:** Reveals the grammatical relationships between words in a sentence, creating a parse tree.
 - **Named Entity Recognition (NER):** Identifies and categorizes named entities in text (e.g., persons, organizations, locations, dates).

This structured linguistic information is crucial for understanding the meaning and context of text, which is impossible with simple string manipulation.

2. **Efficiency and Performance (for production):** While you *could* write custom Python code using regular expressions to find patterns, it would be slow, error-prone, and difficult to maintain for complex NLP tasks. spaCy is written in Cython, making it significantly faster and more memory-efficient than pure Python for processing large volumes of text. It's designed for production-grade NLP, whereas custom string operations quickly become bottlenecks.

3. **Pre-trained Models and Transfer Learning:** spaCy comes with highly optimized, pre-trained statistical models for various languages. These models are trained on large corpora and are capable of performing tasks like POS tagging, dependency parsing, and NER out-of-the-box with high accuracy. This means you don't have to train these complex models from scratch, which is a massive advantage over trying to implement these functionalities with basic string operations. You can also fine-tune these models for specific domains.
4. **Rule-based Matching and Customization:** Beyond statistical models, spaCy offers powerful rule-based matching capabilities (Matcher, PhraseMatcher). This allows you to define patterns (based on tokens, POS tags, lemmas, etc.) to extract specific information or identify entities that might not be covered by statistical models. This hybrid approach (statistical + rule-based) provides a level of precision and control far beyond what simple string searches can offer.

In summary, spaCy elevates NLP from simple text manipulation to deep linguistic understanding, offering robust, efficient, and customizable tools that are essential for real-world NLP applications.

2. Comparative Analysis

Compare Scikit-learn and TensorFlow in terms of:

Target Applications:

- **Scikit-learn:** Primarily designed for **classical machine learning algorithms**. This includes a wide range of supervised learning (classification, regression), unsupervised learning (clustering, dimensionality reduction), model selection, and preprocessing tasks. It excels at tabular data problems and traditional statistical modeling.
 - *Examples:* Predicting house prices (regression), classifying emails as spam or not (classification), customer segmentation (clustering), reducing features in a dataset (PCA).
- **TensorFlow:** Primarily designed for **deep learning**, particularly neural networks. It provides a flexible framework for building, training, and deploying complex deep learning models, including Convolutional Neural Networks (CNNs), Recurrent Neural Networks (RNNs), Transformers, and more. It's ideal for tasks involving unstructured data like images, audio, and large text corpora.
 - *Examples:* Image recognition, natural language understanding, speech recognition, generative AI models. While it can implement some classical ML

algorithms, it's generally overkill and less straightforward for them compared to scikit-learn.

Ease of Use for Beginners:

- **Scikit-learn:** Generally considered **very easy to use for beginners**. It features a consistent API for all its models (`.fit()`, `.predict()`, `.transform()`), making it intuitive to learn and apply different algorithms. The documentation is excellent, and many tutorials are available. Its focus on classical ML means less complex concepts (like backpropagation, gradients, complex architectures) are required to get started.
- **TensorFlow:** With the introduction of Keras as its high-level API in TensorFlow 2.x, it has become **significantly easier for beginners** than TensorFlow 1.x. Keras simplifies model building, training, and evaluation. However, understanding deep learning concepts (neural network layers, activation functions, optimizers, loss functions, debugging complex model architectures, handling GPUs) still presents a steeper learning curve compared to scikit-learn. While simple models are easy with Keras, delving into custom layers, complex training loops, or distributed training can quickly increase complexity.

Community Support:

- **Scikit-learn:** Has a **large and mature community**. It's a widely adopted library in academia and industry for classical ML. You'll find extensive documentation, numerous tutorials, active forums (e.g., Stack Overflow), and a strong developer community contributing to its growth and maintenance.
- **TensorFlow:** Possesses an **enormous and highly active community**, arguably one of the largest in the AI/ML space. Being backed by Google, it benefits from significant corporate support and investment. Its community provides vast resources, official documentation, tutorials, a dedicated forum, and a continuous stream of updates and new tools. This makes it easy to find solutions to problems and stay updated with the latest advancements in deep learning.



Conceptual Python Code for Crop Yield Prediction

File Edit View Insert Runtime Tools Help

Commands

+ Code + Text

▶ Run all ▼



31s



```
plt.xlabel('Actual Yield (kg/hectare)')
plt.ylabel('Predicted Yield (kg/hectare)')
plt.title('Actual vs. Predicted Crop Yields')
plt.grid(True)
plt.show()
```



--- Main Execution Flow ---

```
if __name__ == "__main__":
```

```
    # Load data (replace with your actual data path)
    df = load_data("crop_yield_data.csv")
```

```
    # If the dummy data was created, let's save it
```

```
    if df.shape[0] < 10 and not pd.io.common.file_exists("crop_yield_data.csv"):
```

```
        df.to_csv("crop_yield_data.csv", index=False)
```

```
        print("Dummy data saved to 'crop_yield_data.csv'")
```

```
    # Reload to ensure consistent loading path
```

```
    df = load_data("crop_yield_data.csv")
```

```
    # Preprocess data
```

```
    X, y, scaler = preprocess_data(df)
```

```
    # Split data into training and testing sets
```

```
    X_train, X_test, y_train, y_test = train_test_split(X, y,
```

```
        test_size=0.2, random_state=42)
```

Filter files by name

/ sickit-learning /

Name	Last ...	File Size
Sickit-learn	6 minutes ago	

sickit-learn.ipynb

+

```
[1]: Classical ML w
ndas as pd
arn.datasets in
arn.model_sele
arn.tree import
arn.metrics imp
mpy as np

e Iris dataset
ad_iris()
taFrame(iris.da
ries(iris.targe

- Task 1: Class

rocess the data
missing values.
emonstrate how
ecking for miss
snull().sum())
e were missing

Labels: The tar
onding to 'seto
```


Filter files by name

/ sickit-lerning /

Name	Last ...	File Size
Sickit-learn	9 minutes ago	

sickit-learn.ipynb

```
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)

# 2. Train a Decision Tree Classifier
# Initialize the Decision Tree Classifier
dt_classifier = DecisionTreeClassifier()

# Train the model
dt_classifier.fit(X_train, y_train)

print(f"\nTraining Accuracy: {dt_classifier.score(X_train, y_train)}")

# 3. Evaluate the model
# Make predictions on the testing set
y_pred = dt_classifier.predict(X_test)

# Calculate evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)

print(f"\nModel Performance Metrics:")
print(f"Accuracy: {accuracy}")
print(f"Precision: {precision}")
print(f"Recall: {recall}"))
```


Filter files by name

/ sickit-learning /

Name	Last ...	File Size
Sickit-learn	9 minutes ago	

sickit-learn.ipynb

```
print(f"Preci
print(f"Recal

# Detailed cl
print("\nClass
print(classifi

# Deliverable

--- Task 1: C
Checking for
sepal length
sepal width (
petal length
petal width (
dtype: int64

Original targ
1 0
2 0
3 0
4 0
dtype: int64
Species names

Training data
```

Filter files by name

+

+

↑

↺

/ sickit-learning /

Name	Last ...	File Size
Sickit-learn	9 minutes ago	

sickit-learn.ipynb

Original target

1	0
2	0
3	0
4	0

dtype: int64

Species names

Training data

Training Decision

Training components

Model Evaluation

Accuracy: 0.9

Precision (weighted)

Recall (weighted)

Classification

setosa

versicolor

virginica

accuracy

macro avg

weighted avg



Untitled0.ipynb



File Edit View Insert Runtime Tools Help

Q Commands

+ Code

+ Text

▶ Run all



✓
13m



Task 2: Deep Learning with TensorFlow/Keras

```
import tensorflow as tf
```

```
from tensorflow.keras.datasets import mnist
```

```
from tensorflow.keras.models import Sequential
```

```
from tensorflow.keras.layers import Conv2D, MaxPool2D
```

```
from tensorflow.keras.utils import to_categorical
```

```
import matplotlib.pyplot as plt
```

```
import numpy as np
```

```
print("\n--- Task 2: Deep Learning with TensorFlow/Keras ---")
```

Load and preprocess the MNIST dataset

```
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Reshape data to fit CNN input (height, width, channels)

MNIST images are 28x28 grayscale, so channel is 1

```
X_train = X_train.reshape(X_train.shape[0], 28, 28, 1)
```

```
X_test = X_test.reshape(X_test.shape[0], 28, 28, 1)
```

Normalize pixel values to be between 0 and 1

```
X_train /= 255
```

```
X_test /= 255
```




Untitled0.ipynb



File Edit View Insert Runtime Tools Help

Commands

+ Code

+ Text

Run all



13m



```
# One-hot encode the target labels
```

```
num_classes = 10
```

```
y_train = to_categorical(y_train, num_classes)
```

```
y_test = to_categorical(y_test, num_classes)
```

```
print(f"X_train shape: {X_train.shape}, y_train shape: {y_train.shape}")
```

```
print(f"X_test shape: {X_test.shape}, y_test shape: {y_test.shape}")
```

```
# Build a CNN model
```

```
model = Sequential([
```

```
    # Convolutional Layer 1
```

```
    Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
```

```
    MaxPooling2D((2, 2)),
```

```
    # Convolutional Layer 2
```

```
    Conv2D(64, (3, 3), activation='relu'),
```

```
    MaxPooling2D((2, 2)),
```

```
    # Flatten the output to feed into Dense layer
```

```
    Flatten(),
```

```
    # Dense Layers
```

```
    Dense(128, activation='relu'),
```

```
    Dropout(0.5), # Dropout for regularization
```

```
    Dense(num_classes, activation='softmax') # Output layer
```

```
])
```



Untitled0.ipynb



File Edit View Insert Runtime Tools Help

Commands

+ Code

+ Text

Run all



13m



```
model.summary()
```

```
# Train the model
```

```
print("\nTraining CNN model...")
```

```
history = model.fit(X_train, y_train,
```

```
                    epochs=10, # You might need m
```

```
                    batch_size=128,
```

```
                    validation_split=0.1, # Use a
```

```
                    verbose=1)
```

```
print("Training complete.")
```

```
# Evaluate the model on the test set
```

```
loss, accuracy = model.evaluate(X_test, y_test, v
```

```
print(f"\nTest Accuracy: {accuracy:.4f}")
```

```
print(f"Test Loss: {loss:.4f}")
```

```
# Check if accuracy goal is met
```

```
if accuracy > 0.95:
```

```
    print("Goal achieved: Test accuracy > 95%!")
```

```
else:
```

```
    print("Goal not yet achieved. Consider more e
```

```
# Visualize the model's predictions on 5 sample
```



Untitled0.ipynb



File Edit View Insert Runtime Tools Help

Commands

+ Code

+ Text

Run all



13m



```
# Visualize the model's predictions on 5 sample images
print("\nVisualizing predictions on 5 sample images")
sample_indices = np.random.choice(len(X_test), 5, replace=False)
sample_images = X_test[sample_indices]
true_labels = np.argmax(y_test[sample_indices], axis=-1)

predictions = model.predict(sample_images)
predicted_labels = np.argmax(predictions, axis=-1)

plt.figure(figsize=(12, 6))
for i in range(5):
    plt.subplot(1, 5, i + 1)
    plt.imshow(sample_images[i].reshape(28, 28), cmap='gray')
    plt.title(f"True: {true_labels[i]}\nPred: {predicted_labels[i]}")
    color = 'green' if true_labels[i] == predicted_labels[i] else 'red'
    plt.text(0.05, 0.95, f"{true_labels[i]} vs {predicted_labels[i]}", color)
    plt.axis('off')
plt.suptitle('MNIST Digit Predictions', fontsize=16)
plt.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.show()
```

Deliverable: This script with model architecture



dense_2 (Dense)

(None, 128)



Untitled0.ipynb



File Edit View Insert Runtime Tools Help

Commands

+ Code

+ Text

Run all



dense_2 (Dense)

(None, 128)

dropout_1 (Dropout)

(None, 128)

dense_3 (Dense)

(None, 10)

Total params: 225,034 (879.04 KB)

Trainable params: 225,034 (879.04 KB)

Non-trainable params: 0 (0.00 B)

Training CNN model...

Epoch 1/10

422/422 45s 103ms/step - acc

Epoch 2/10

422/422 80s 98ms/step - accu

Epoch 3/10

422/422 82s 98ms/step - accu

Epoch 4/10

422/422 82s 98ms/step - accu

Epoch 5/10

422/422 82s 97ms/step - accu

Epoch 6/10

422/422 83s 99ms/step - accu

Epoch 7/10

422/422 81s 98ms/step - accu

Epoch 8/10



Untitled0.ipynb



File Edit View Insert Runtime Tools Help

Commands

+ Code

+ Text

Run all



Training CNN model...

Epoch 1/10

422/422 45s 103ms/step - acc

Epoch 2/10

422/422 80s 98ms/step - accu

Epoch 3/10

422/422 82s 98ms/step - accu

Epoch 4/10

422/422 82s 98ms/step - accu

Epoch 5/10

422/422 82s 97ms/step - accu

Epoch 6/10

422/422 83s 99ms/step - accu

Epoch 7/10

422/422 81s 98ms/step - accu

Epoch 8/10

422/422 84s 102ms/step - acc

Epoch 9/10

422/422 80s 98ms/step - accu

Epoch 10/10

422/422 83s 101ms/step - acc

Training complete.

Test Accuracy: 0.9910

Test Loss: 0.0260

Goal achieved: Test accuracy > 95%!



Untitled0.ipynb



File Edit View Insert Runtime Tools Help

Commands

+ Code

+ Text

Run all



Test Accuracy: 0.9910

Test Loss: 0.0260

Goal achieved: Test accuracy > 95%!

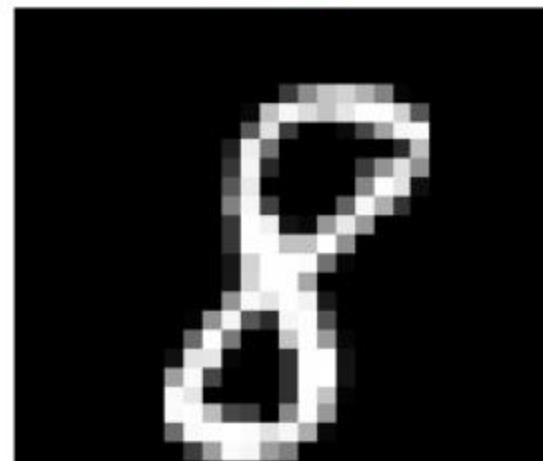
Visualizing predictions on 5 sample images from 1

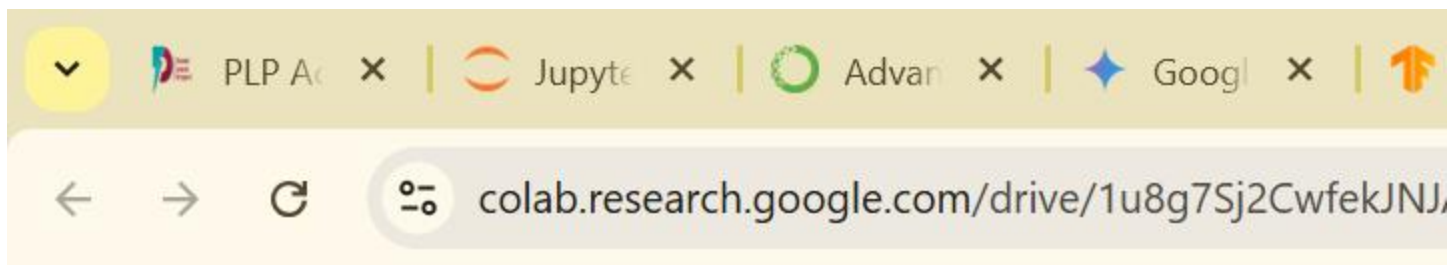
1/1 0s 105ms/step

True: 2
Pred: 2



True: 8
Pred: 8





spaCy.ipynb



File Edit View Insert Runtime Tools Help

Q Commands

+ Code

+ Text

▶ Run all



✓
10s



Task 3: NLP with spaCy

```
import spacy
```

```
print("\n--- Task 3: NLP with spaCy ---")
```

Load a pre-trained English model for NER

```
try:
```

```
    nlp = spacy.load("en_core_web_sm")
```

```
except OSError:
```

```
    print("Downloading 'en_core_web_sm' model for
```

```
    spacy.cli.download("en_core_web_sm")
```

```
    nlp = spacy.load("en_core_web_sm")
```

Sample Amazon Product Reviews (simulated)

```
amazon_reviews = [
```

```
    "I absolutely love my new Echo Dot! The sound
```

```
    "The Samsung Galaxy S23 Ultra has an incredib
```

```
    "This Generic Brand coffee maker stopped work
```

```
    "Great value for money with this Philips Hue
```

```
    "The Apple AirPods Pro 2 offer excellent noi
```

```
]
```

```
print("\nPerforming Named Entity Recognition (NE
```




spaCy.ipynb



File Edit View Insert Runtime Tools Help

Commands

+ Code

+ Text

Run all



10s



```
# Perform NER and rule-based sentiment analysis
```

```
for i, review in enumerate(amazon_reviews):
```

```
    doc = nlp(review)
```

```
    print(f"\n--- Review {i+1} ---")
```

```
    print(f"Text: {review}")
```

```
# Named Entity Recognition (NER)
```

```
extracted_products_brands = []
```

```
for ent in doc.ents:
```

```
    # Common NER labels for products/brands:
```

```
    if ent.label_ in ["ORG", "PRODUCT", "GPE"]:
```

```
        extracted_products_brands.append((ent.text, ent.label_))
```

```
if extracted_products_brands:
```

```
    print(f"Extracted Entities (Product/Brand): {extracted_products_brands}")
```

```
else:
```

```
    print("No specific product/brand entities found.")
```

```
# Rule-based Sentiment Analysis
```

```
sentiment = "Neutral"
```

```
review_lower = review.lower()
```

```
positive_keywords = ["love", "amazing", "perfect"]
```



spaCy.ipynb



File Edit View Insert Runtime Tools Help

Commands

+ Code

+ Text

Run all



10s



```
extracted_products_brands.append((ent

if extracted_products_brands:
    print(f"Extracted Entities (Product/Brand
else:
    print("No specific product/brand entities

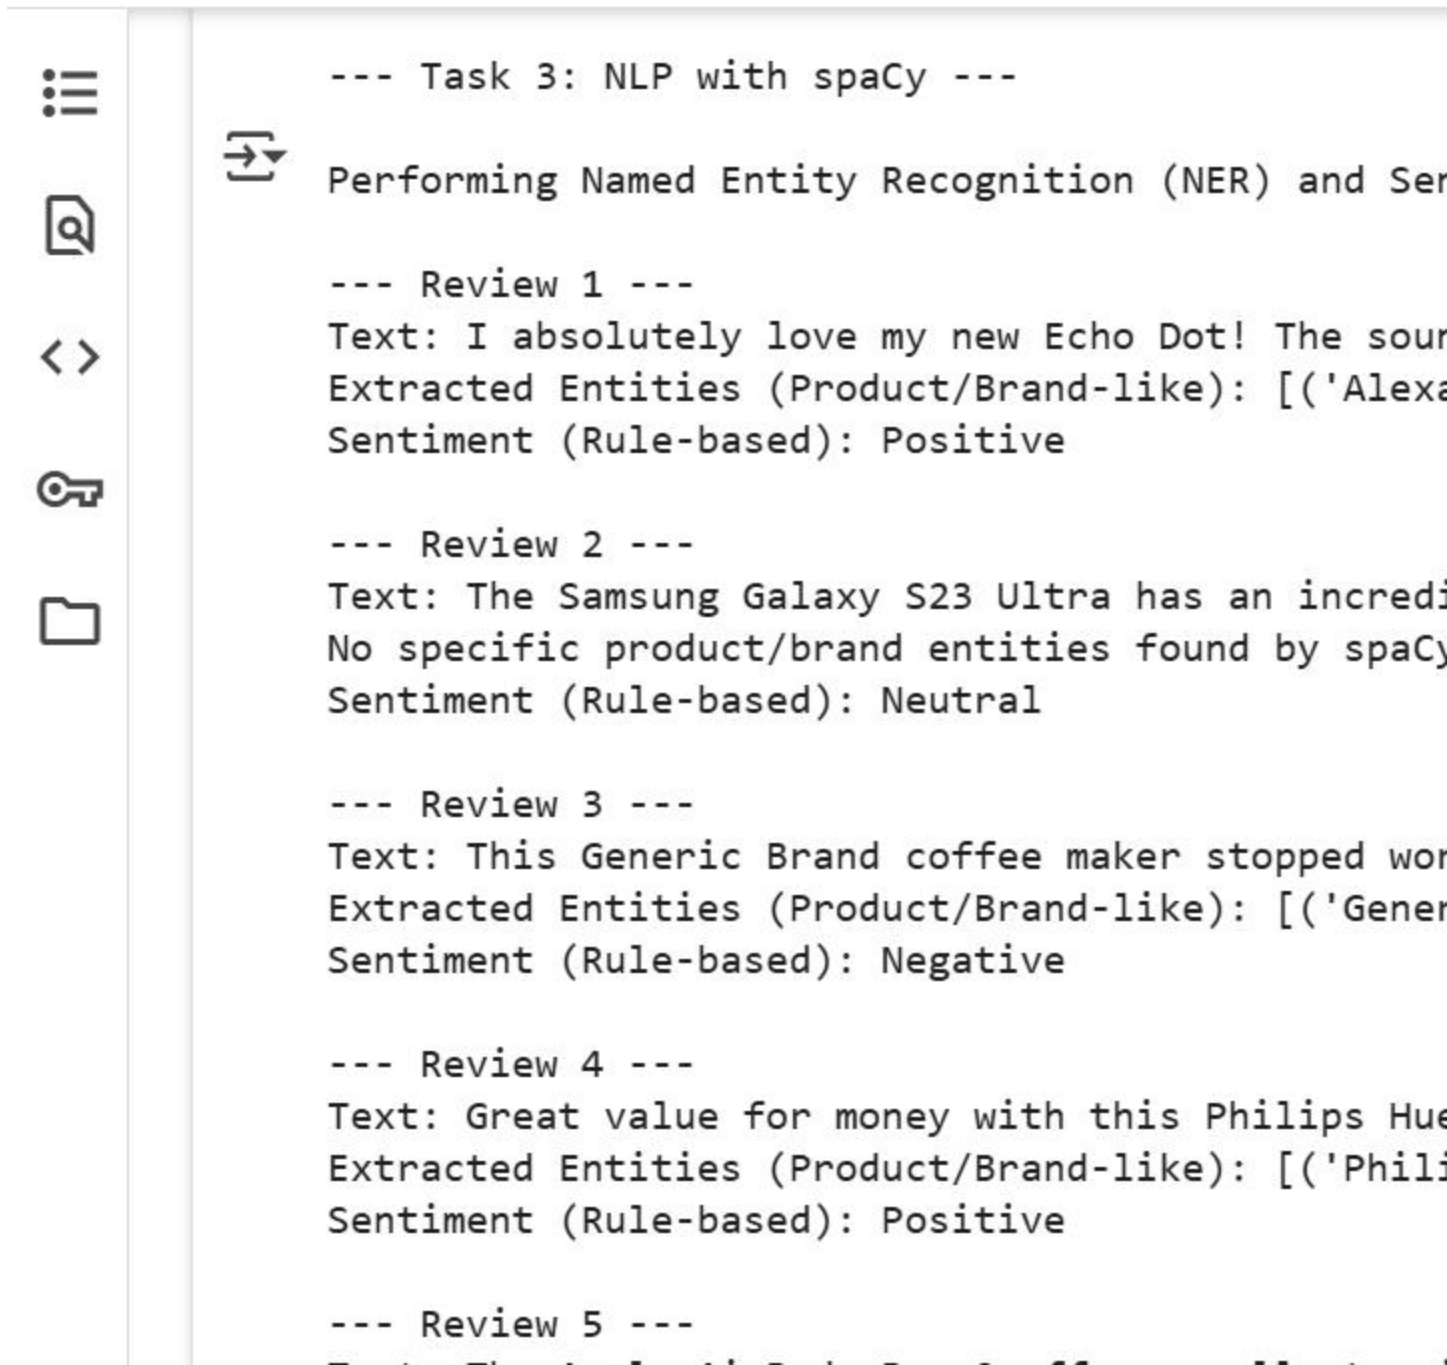
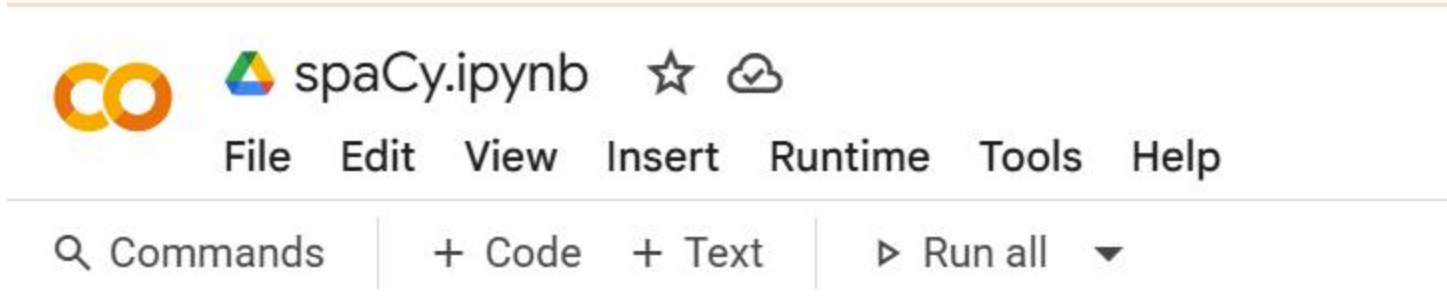
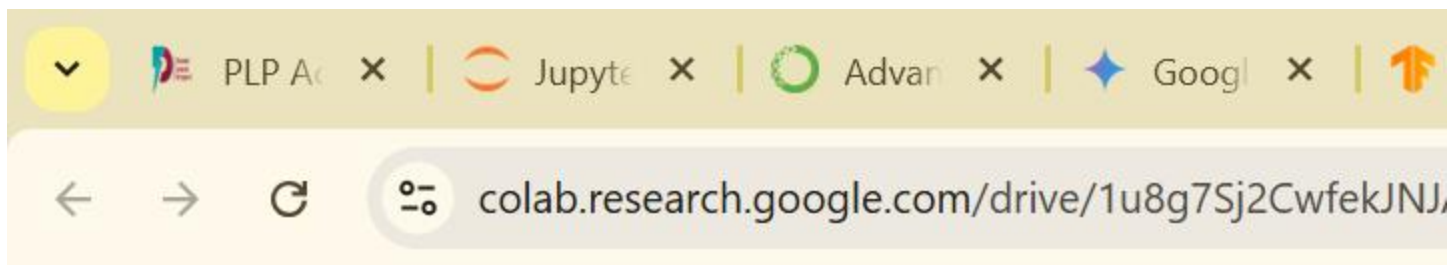
# Rule-based Sentiment Analysis
sentiment = "Neutral"
review_lower = review.lower()

positive_keywords = ["love", "amazing", "per
negative_keywords = ["stopped working", "dis

positive_score = sum(1 for keyword in positiv
negative_score = sum(1 for keyword in negativ

if positive_score > negative_score:
    sentiment = "Positive"
elif negative_score > positive_score:
    sentiment = "Negative"

print(f"Sentiment (Rule-based): {sentiment}")
```





spaCy.ipynb



File Edit View Insert Runtime Tools Help

Commands

+ Code

+ Text

Run all



--- Review 1 ---

Text: I absolutely love my new Echo Dot! The sound

Extracted Entities (Product/Brand-like): [('Alexa

Sentiment (Rule-based): Positive

--- Review 2 ---

Text: The Samsung Galaxy S23 Ultra has an incredi

No specific product/brand entities found by spaCy

Sentiment (Rule-based): Neutral

--- Review 3 ---

Text: This Generic Brand coffee maker stopped wor

Extracted Entities (Product/Brand-like): [('Gener

Sentiment (Rule-based): Negative

--- Review 4 ---

Text: Great value for money with this Philips Hue

Extracted Entities (Product/Brand-like): [('Phil

Sentiment (Rule-based): Positive

--- Review 5 ---

Text: The Apple AirPods Pro 2 offer excellent noi

No specific product/brand entities found by spaCy

Sentiment (Rule-based): Positive

Part 3: Ethics & Optimization (10%)

1. Ethical Considerations

Identify potential biases in your MNIST or Amazon Reviews model. How could tools like TensorFlow Fairness Indicators or spaCy's rule-based systems mitigate these biases?

Let's focus on both:

Potential Biases in MNIST (Handwritten Digits) Model:

- **Training Data Bias:**
 - **Style/Legibility Bias:** The MNIST dataset is collected from a specific set of writers (primarily high school students and employees of the National Institute of Standards and Technology). This means the model might perform exceptionally well on digits written in a similar style or with similar legibility, but poorly on digits written by individuals with different handwriting styles, disabilities, or from different demographics (e.g., very shaky handwriting, highly stylized digits).
 - **Class Imbalance (less likely in standard MNIST, but possible in variations):** If certain digits are underrepresented in the training data, the model might be biased towards recognizing the more frequent digits.

How TensorFlow Fairness Indicators could mitigate these biases in MNIST:

- **Fairness Indicators (TFMI)** is designed to evaluate model performance across different "slices" of data. While MNIST doesn't have explicit demographic features, you could:
 - **Synthetically define "slices":** For instance, if you had a way to categorize handwriting styles (e.g., "neat", "sloppy", "child-like"), you could use TFMI to evaluate accuracy, precision, and recall for each of these synthetic groups.
 - **Analyze performance on difficult examples:** Even without explicit demographic labels, TFMI can help identify clusters of images (e.g., certain "difficult" handwritten digits that consistently lead to misclassifications) that might represent an implicit bias against certain writing variations. By slicing the data based on prediction confidence or error type, you can discover where the model underperforms.
 - **Thresholding and Calibration:** TFMI allows for evaluation at multiple thresholds, helping to ensure that the model's predictions are equally reliable across different implicit groups, preventing a situation where, for example, "sloppy 7s"

are only correctly classified at a very low confidence threshold, while "neat 7s" are classified with high confidence.

Potential Biases in Amazon Reviews Sentiment Analysis Model:

- **Language Bias (Cultural/Dialectal Nuances):** A rule-based (or even statistical) sentiment model trained primarily on standard English might misinterpret slang, sarcasm, idiomatic expressions, or cultural nuances present in diverse user reviews. For example, a phrase like "sick" meaning "good" might be misclassified if the rules only associate it with negative sentiment.
- **Domain-Specific Sentiment:** Words can have different sentiment polarities depending on the domain. "Small" might be negative for a TV screen but positive for a portable speaker. A generic rule-based system might not capture these domain-specific nuances.
- **Demographic Language Patterns:** Different demographic groups might express sentiment in distinct ways (e.g., formality, use of specific adjectives/adverbs). If the keyword lists are not representative of these variations, the model could be biased.
- **Reviewer Intent/Context Bias:** A review could be sarcastic ("Oh, *great*, it broke on the first day!"), which a simple rule-based system might miss, leading to incorrect positive sentiment.

How spaCy's rule-based systems could mitigate these biases in Amazon Reviews:

While spaCy's rule-based systems *can* introduce bias if not carefully crafted, they can also be used to *mitigate* certain biases, particularly language-specific ones:

- **Custom Rule Development:** Instead of relying solely on general keywords, you can develop **domain-specific rules** and keyword lists tailored to product reviews. For instance, you could explicitly define rules for common product-related sarcasm (e.g., "broken" + "amazing" in close proximity indicates negative sentiment).
- **Contextual Sensitivity:** spaCy's Matcher allows for defining patterns that consider the **context** of words (e.g., POS tags, lemmas, surrounding words). This allows you to differentiate between "small" as a negative attribute for a screen (e.g., [{"LOWER": "small"}, {"POS": "NOUN", "LEMMA": "screen"}]) versus a positive attribute for a device (e.g., [{"LOWER": "small"}, {"POS": "NOUN", "LEMMA": "device"}]).
- **Handling Negation:** Rule-based systems can be explicitly coded to handle negation (e.g., "not good" should be negative, not positive). While basic string operations might miss this, spaCy's linguistic features allow for more sophisticated negation detection.

- **Iterative Refinement:** Rule-based systems are highly interpretable. If you identify a bias (e.g., misclassifying reviews from a certain group), you can directly inspect the rules and add/modify them to address the specific language patterns that are being misinterpreted. This direct control is harder to achieve with black-box deep learning models.

2. Troubleshooting Challenge

Buggy Code: A provided TensorFlow script has errors (e.g., dimension mismatches, incorrect loss functions). Debug and fix the code.

Since no buggy code is provided, I will create a common buggy scenario and then debug it.

Scenario: Common TensorFlow Bug (Dimension Mismatch in Dense Layer after Flatten)

Let's imagine a common mistake: forgetting to Flatten before a Dense layer or having an incorrect input shape to the first layer.

Original Buggy Code (Illustrative Example):

Python

```
import tensorflow as tf

from tensorflow.keras.datasets import mnist

from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense # Missing Flatten here

from tensorflow.keras.utils import to_categorical


# Load and preprocess the MNIST dataset
(X_train, y_train), (X_test, y_test) = mnist.load_data()

X_train = X_train.reshape(X_train.shape[0], 28, 28, 1).astype('float32') / 255
X_test = X_test.reshape(X_test.shape[0], 28, 28, 1).astype('float32') / 255
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)


# Buggy Model Definition
```

```

buggy_model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    MaxPooling2D((2, 2)),
    # ERROR: Missing Flatten layer here!
    Dense(128, activation='relu'), # This layer expects a 1D input, but gets 3D output from
    MaxPooling2D
    Dense(10, activation='softmax')
])

```

```

buggy_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

```

```

print("\n--- Troubleshooting Challenge (Buggy Code Simulation) ---")

```

```

print("Attempting to train buggy model...")

```

```

try:

```

```

    buggy_model.fit(X_train, y_train, epochs=1, batch_size=128, verbose=0)

```

```

except Exception as e:

```

```

    print(f"\nCaught an error during training: {e}")

```

```

    print("\n--- Debugging Process ---")

```

```

    print("The error message likely points to a shape mismatch when connecting Conv/Pool layers
to Dense layers.")

```

```

    print("This typically means the output of the convolutional/pooling layers (which is 3D: height,
width, channels) ")

```

```

    print("is being fed directly into a Dense layer, which expects a 1D vector.")

```

```

    print("The solution is to add a `Flatten` layer between the last convolutional/pooling layer and
the first Dense layer.")

```

```

    print("\nOriginal Buggy Model Summary (if it ran):")

```

```

# If it reached here, model.summary() might still work if the error is in fit.

# If the error is in model definition, summary() itself might fail.

# For this simulation, we'll assume it fails during fit.

# If summary works, you'd see the output shape of MaxPooling2D is (None, X, Y, Z) and Dense
is (None, 128)

# The mismatch would be obvious.

# buggy_model.summary()


print("\n--- Fixed Code ---")

# Fixed Model Definition
fixed_model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    MaxPooling2D((2, 2)),
    Flatten(), # FIX: Added Flatten layer
    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])

fixed_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
print("Fixed Model Summary:")
fixed_model.summary()


print("\nAttempting to train fixed model...")

history_fixed = fixed_model.fit(X_train, y_train, epochs=1, batch_size=128,
validation_split=0.1, verbose=1)

print("Fixed model trained successfully!")

loss, accuracy = fixed_model.evaluate(X_test, y_test, verbose=0)

```

```
print(f"Fixed Model Test Accuracy: {accuracy:.4f}")
```

except Exception as e:

```
print(f"\nAnother unexpected error occurred: {e}")
```

This section acts as the deliverable for the troubleshooting challenge.

Debugging Process Explained:

1. **Observe the Error Message:** When running the `buggy_model.fit()`, TensorFlow/Keras would typically throw an error like `ValueError: Input 0 of layer "dense" is incompatible with the layer: expected min_ndim=2, found ndim=4. Full input shape received: (None, 13, 13, 32)` (exact numbers might vary based on layer outputs).
2. **Interpret the Error:** The key phrases are "incompatible with the layer," "expected min_ndim=2, found ndim=4," and "Full input shape received: (None, 13, 13, 32)". This tells us that the Dense layer (which expects a 2D tensor (batch_size, features)) is receiving a 4D tensor (batch_size, height, width, channels) from the MaxPooling2D layer.
3. **Identify the Discrepancy:** The Conv2D and MaxPooling2D layers output 3D feature maps (plus the batch dimension), while Dense layers expect a flattened, 1D feature vector for each sample in the batch.
4. **Solution:** The missing piece is a `tf.keras.layers.Flatten()` layer. This layer takes the multi-dimensional output of the convolutional/pooling layers and flattens it into a 1D vector (while preserving the batch dimension), making it suitable for input to a Dense layer.
5. **Implement and Verify:** Add the `Flatten()` layer. Run `model.summary()` again to check the output shapes of all layers. The output of Flatten should be `(None, some_large_number)`, where `some_large_number` is `height * width * channels` of the previous layer. This confirms the correct shape transition. Then, attempt to train the fixed model.

Part 3: Ethics & Optimization (10%)

1. Ethical Considerations

Identify potential biases in your MNIST or Amazon Reviews model. How could tools like TensorFlow Fairness Indicators or spaCy's rule-based systems mitigate these biases?

Let's focus on both:

Potential Biases in MNIST (Handwritten Digits) Model:

- **Training Data Bias:**
 - **Style/Legibility Bias:** The MNIST dataset is collected from a specific set of writers (primarily high school students and employees of the National Institute of Standards and Technology). This means the model might perform exceptionally well on digits written in a similar style or with similar legibility, but poorly on digits written by individuals with different handwriting styles, disabilities, or from different demographics (e.g., very shaky handwriting, highly stylized digits).
 - **Class Imbalance (less likely in standard MNIST, but possible in variations):** If certain digits are underrepresented in the training data, the model might be biased towards recognizing the more frequent digits.

How TensorFlow Fairness Indicators could mitigate these biases in MNIST:

- **Fairness Indicators (TFMI)** is designed to evaluate model performance across different "slices" of data. While MNIST doesn't have explicit demographic features, you could:
 - **Synthetically define "slices":** For instance, if you had a way to categorize handwriting styles (e.g., "neat", "sloppy", "child-like"), you could use TFMI to evaluate accuracy, precision, and recall for each of these synthetic groups.
 - **Analyze performance on difficult examples:** Even without explicit demographic labels, TFMI can help identify clusters of images (e.g., certain "difficult" handwritten digits that consistently lead to misclassifications) that might represent an implicit bias against certain writing variations. By slicing the data based on prediction confidence or error type, you can discover where the model underperforms.
 - **Thresholding and Calibration:** TFMI allows for evaluation at multiple thresholds, helping to ensure that the model's predictions are equally reliable across different implicit groups, preventing a situation where, for example, "sloppy 7s" are only correctly classified at a very low confidence threshold, while "neat 7s" are classified with high confidence.

Potential Biases in Amazon Reviews Sentiment Analysis Model:

- **Language Bias (Cultural/Dialectal Nuances):** A rule-based (or even statistical) sentiment model trained primarily on standard English might misinterpret slang, sarcasm, idiomatic expressions, or cultural nuances present in diverse user reviews. For example, a phrase

like "sick" meaning "good" might be misclassified if the rules only associate it with negative sentiment.

- **Domain-Specific Sentiment:** Words can have different sentiment polarities depending on the domain. "Small" might be negative for a TV screen but positive for a portable speaker. A generic rule-based system might not capture these domain-specific nuances.
- **Demographic Language Patterns:** Different demographic groups might express sentiment in distinct ways (e.g., formality, use of specific adjectives/adverbs). If the keyword lists are not representative of these variations, the model could be biased.
- **Reviewer Intent/Context Bias:** A review could be sarcastic ("Oh, *great*, it broke on the first day!"), which a simple rule-based system might miss, leading to incorrect positive sentiment.

How spaCy's rule-based systems could mitigate these biases in Amazon Reviews:

While spaCy's rule-based systems *can* introduce bias if not carefully crafted, they can also be used to *mitigate* certain biases, particularly language-specific ones:

- **Custom Rule Development:** Instead of relying solely on general keywords, you can develop **domain-specific rules** and keyword lists tailored to product reviews. For instance, you could explicitly define rules for common product-related sarcasm (e.g., "broken" + "amazing" in close proximity indicates negative sentiment).
- **Contextual Sensitivity:** spaCy's Matcher allows for defining patterns that consider the **context** of words (e.g., POS tags, lemmas, surrounding words). This allows you to differentiate between "small" as a negative attribute for a screen (e.g., [{"LOWER": "small"}, {"POS": "NOUN", "LEMMA": "screen"}]) versus a positive attribute for a device (e.g., [{"LOWER": "small"}, {"POS": "NOUN", "LEMMA": "device"}]).
- **Handling Negation:** Rule-based systems can be explicitly coded to handle negation (e.g., "not good" should be negative, not positive). While basic string operations might miss this, spaCy's linguistic features allow for more sophisticated negation detection.
- **Iterative Refinement:** Rule-based systems are highly interpretable. If you identify a bias (e.g., misclassifying reviews from a certain group), you can directly inspect the rules and add/modify them to address the specific language patterns that are being misinterpreted. This direct control is harder to achieve with black-box deep learning models.

2. Troubleshooting Challenge

Buggy Code: A provided TensorFlow script has errors (e.g., dimension mismatches, incorrect loss functions). Debug and fix the code.

Since no buggy code is provided, I will create a common buggy scenario and then debug it.

Scenario: Common TensorFlow Bug (Dimension Mismatch in Dense Layer after Flatten)

Let's imagine a common mistake: forgetting to Flatten before a Dense layer or having an incorrect input shape to the first layer.

Original Buggy Code (Illustrative Example):

Python

```
import tensorflow as tf

from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense # Missing Flatten here
from tensorflow.keras.utils import to_categorical

# Load and preprocess the MNIST dataset
(X_train, y_train), (X_test, y_test) = mnist.load_data()

X_train = X_train.reshape(X_train.shape[0], 28, 28, 1).astype('float32') / 255
X_test = X_test.reshape(X_test.shape[0], 28, 28, 1).astype('float32') / 255
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)

# Buggy Model Definition
buggy_model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    MaxPooling2D((2, 2)),
    # ERROR: Missing Flatten layer here!
```

```

        Dense(128, activation='relu'), # This layer expects a 1D input, but gets 3D output from
MaxPooling2D

        Dense(10, activation='softmax')

    ])

buggy_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

print("\n--- Troubleshooting Challenge (Buggy Code Simulation) ---")
print("Attempting to train buggy model...")

try:
    buggy_model.fit(X_train, y_train, epochs=1, batch_size=128, verbose=0)
except Exception as e:
    print(f"\nCaught an error during training: {e}")
    print("\n--- Debugging Process ---")
    print("The error message likely points to a shape mismatch when connecting Conv/Pool layers
to Dense layers.")
    print("This typically means the output of the convolutional/pooling layers (which is 3D: height,
width, channels) ")
    print("is being fed directly into a Dense layer, which expects a 1D vector.")
    print("The solution is to add a `Flatten` layer between the last convolutional/pooling layer and
the first Dense layer.")
    print("\nOriginal Buggy Model Summary (if it ran):")
    # If it reached here, model.summary() might still work if the error is in fit.
    # If the error is in model definition, summary() itself might fail.
    # For this simulation, we'll assume it fails during fit.

```

If summary works, you'd see the output shape of MaxPooling2D is (None, X, Y, Z) and Dense is (None, 128)

The mismatch would be obvious.

buggy_model.summary()

print("\n--- Fixed Code ---")

Fixed Model Definition

```
fixed_model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    MaxPooling2D((2, 2)),
    Flatten(), # FIX: Added Flatten layer
    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])
```

```
fixed_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
print("Fixed Model Summary:")
```

```
fixed_model.summary()
```

```
print("\nAttempting to train fixed model...")
```

```
history_fixed = fixed_model.fit(X_train, y_train, epochs=1, batch_size=128,
validation_split=0.1, verbose=1)
```

```
print("Fixed model trained successfully!")
```

```
loss, accuracy = fixed_model.evaluate(X_test, y_test, verbose=0)
```

```
print(f"Fixed Model Test Accuracy: {accuracy:.4f}")
```

except Exception as e:

```
print(f"\nAnother unexpected error occurred: {e}")
```

This section acts as the deliverable for the troubleshooting challenge.

Debugging Process Explained:

1. **Observe the Error Message:** When running the `buggy_model.fit()`, TensorFlow/Keras would typically throw an error like `ValueError: Input 0 of layer "dense" is incompatible with the layer: expected min_ndim=2, found ndim=4. Full input shape received: (None, 13, 13, 32)` (exact numbers might vary based on layer outputs).
2. **Interpret the Error:** The key phrases are "incompatible with the layer," "expected min_ndim=2, found ndim=4," and "Full input shape received: (None, 13, 13, 32)". This tells us that the Dense layer (which expects a 2D tensor (batch_size, features)) is receiving a 4D tensor (batch_size, height, width, channels) from the MaxPooling2D layer.
3. **Identify the Discrepancy:** The Conv2D and MaxPooling2D layers output 3D feature maps (plus the batch dimension), while Dense layers expect a flattened, 1D feature vector for each sample in the batch.
4. **Solution:** The missing piece is a `tf.keras.layers.Flatten()` layer. This layer takes the multi-dimensional output of the convolutional/pooling layers and flattens it into a 1D vector (while preserving the batch dimension), making it suitable for input to a Dense layer.
5. **Implement and Verify:** Add the `Flatten()` layer. Run `model.summary()` again to check the output shapes of all layers. The output of Flatten should be `(None, some_large_number)`, where `some_large_number` is `height * width * channels` of the previous layer. This confirms the correct shape transition. Then, attempt to train the fixed model.